

Performance and security exploration of aggressor-focussed row-migration using row swap approach and it's enhancements

Hena Naaz

Abstract—With increasing DRAM scaling, the threat of row hammer attacks on memory systems are increasing rapidly. In this project, we span the implementation of aggressor-focused like RRS in a more performance efficient model. Since the security, performance overhead, and complexity in design remains a concern, in this paper, we present an optimized implementation of RRS/SSRS that achieves higher performance efficiency compared to previous work and presents a secure system. The goal of the project is to optimize the row-migration implementation to reduce overheads. This work contributes to the ongoing effort to develop scalable and secure techniques for mitigating row hammer attacks in DRAM systems.



1 INTRODUCTION

Row hammer attacks pose a significant threat to the security and reliability of DRAM systems, and as DRAM scaling continues, the risk of these attacks is only increasing with smaller row hammer thresholds. Several mitigation techniques have been proposed to address this issue, but these techniques often come with performance overhead, which can limit their practicality.

Due to the nature of attacks on victim-focused mitigation, the focus has been shifted on aggressor-focused mitigation. One promising technique for mitigating Row Hammer attacks is the extension of Randomized Row Swap (RRS)/Scalable and Secure Row Swap (SSRS) approach with a more effective tracker, which can achieve a more efficient design solution while also offering promise for future scalability. However, the performance implications of this technique have yet to be fully addressed.

In this paper, we aim to address these issues by investigating the RRS technique and exploring ways to optimize its performance. Specifically, we have implemented RRS with the baseline misra-gries tracker to compare with baseline (victim refresh) and evaluated its performance using a range of benchmarks and workloads. We have also explored different approaches for improving the performance of RRS, including extending the design with SRAM based trackers and investigating design enhancements for reducing memory overheads across swap handling. [2] [3]

The project aims to contribute to the ongoing effort to develop more secure and efficient aggressor-focused techniques for mitigating Row Hammer attacks in DRAM systems.

2 BACKGROUND

Row hammer attacks are a well-known and growing threat to the security and reliability of DRAM systems. Various mitigation techniques have been proposed to address this issue, including ECC memory, Target Row Refresh (TRR), Row Clone, and Randomized Row Swap. While these techniques have shown promise, they also have limitations such as high cost, limited scalability, and limited effectiveness against sophisticated row hammer attacks. [3]

The RRS/SSRS row-swap technique was proposed as a means of addressing these limitations. This technique

utilizes a hybrid row swap algorithm to randomly swap rows in DRAM, making it difficult for attackers to predict which rows will be affected by a row hammer attack. [2] Despite these benefits, however, it also has limitations in the reliance on probabilistic nature of mitigation and a software-based security algorithm, which may not be accurate or effective in all scenarios.

Another limitation of RRS is the potential impact on memory access latency and bandwidth due to the additional overheads introduced by the hybrid row swap algorithm that needs to maintain history of swap for un-swaps. This overhead can reduce the overall performance of the system and limit its practicality in certain scenarios. Therefore, optimizing the algorithm and its implementation is crucial to minimize these overheads and improve the overall performance of the SSRS technique.

In this paper, we aim to build upon the existing research on RRS/SSRS by exploring ways to improve its performance and security. There are 3 approaches that have been practiced for this, targeting the detection, mitigation, and performance of the design. These are described in the design section of the paper.

3 METHODOLOGY

There are limitations of the discussed row-swap techniques, that includes the potential impact on memory access latency and bandwidth due to the additional overheads introduced by the hybrid row swap and un-swap operations. To address these, the project could explore the optimization of the algorithm and its implementation.

This involved techniques such as tracker modification, data structure optimizations, and smarter scheduling of memory operations to minimize the overheads, thereby, improving the performance of the technique. Furthermore, the safety aspect of the design can be discussed for better row-migration schemes.

The proposed design section will explain in detail the following three hypotheses:

- 1) RRS/SSRS design will improve in performance and overhead with an advanced tracker like Hydra with SRAM based storage as hybrid tracker. [1]
- 2) RRS/SSRS design will improve by replacing the row indirection table that gets exploded and cause

memory overhead with a tracker map that keeps pointers for the swap for un-swaps.

- 3) The design can be made more secure by replacing the row-swap randomization approach with an isolation based approach, keeping a pointer to the migrated rows. [4]

The scope of this project has tested these hypothesis in a simulation environment to study the memory overheads and performance implications of the proposed methodology. In order to perform these experiments, the following configuration of Simplesim simulator was used:

Criterion	Used methodology
Simulator	Simple Sim (CS7292 Lab1 simulator)
Benchmarks	(SPEC) bzip, bwaves, Gems, cactus, mcf, zeusmp
Configuration	4 CPU OoO cores and multi-level set-associative caches
Memory config	1 Channel * 1 rank * 16 banks Memory Size: 16 GB, Row Size: 8 KB
Metrics	Number of migrations, total sys cycles, mitigations, total activations, number of entries in tables, memory overhead.

Fig. 1. The methodology used to perform the stated experiments.

4 DESIGN EXPLORATION

4.1 Overview

The proposed design extends the Randomized Row-Swap (RRS) technique by replacing the spill-count based Misra-Gries tracker with more reliable tracker. It also tries to reduce the memory overhead of storing the tracker and indirection table by exploring integrated structures.

There are certain aspects of the RRS design that are of particular interest to the project. The row indirection table is used to map the addresses of DRAM rows to swapped addresses, which are then used to perform row swaps. The optimization and better management of this table to reduce the overhead and improve the overall performance of the technique is another track explored. [2]




Fig. 2. Row Indirection Table (real and mirrored) provides indirection to the rows involved in the swap operations in RRS/SSRS.

When a row swap is performed, the row-indirection table is updated for the corresponding bank. This information is used to determine un-swaps when it is necessary to perform a full refresh of the DRAM bank. The project plans to optimize the management of these data structures to reduce the overhead and improve the scalability of the technique. Further, the goal is to explore row-migration reduction and security improvement leveraging other aggressor-focused mitigation schemes.

4.2 SRAM based Row Hammer tracker

It is known that the Hydra tracker, which is a hybrid tracker between misra-gries (used in RRS/SSRS) and CRA tracker, is a more efficient and accurate tracker that uses a set of counters to keep track of the number of times each row has been accessed, by creating a cache for tracking. This allows the system to detect and mitigate Row Hammer attacks more effectively.

Graphene Tracker using Misra Gries might have slightly lower accuracy metrics due to its simpler tracking mechanism, leading to potential under-detection or over-detection of Row Hammer vulnerable rows. Hydra achieves higher detection accuracy compared to Graphene due to its hybrid approach, combining algorithmic tracking for initial identification of aggressor rows and exact tracking. [1]

Hydra leverages SRAM-based structure that tracks aggregated counts at the granularity of a group of rows and a per row tracker stored in the DRAM-array which can track arbitrary number of rows. It was observed that with Hydra tracker, the main memory overhead for tracker structure has seen a reduction in magnitude. However, the overall memory overhead doesn't benefit from this change due to the larger overhead coming from the row-indirection table.

Although the perspective of SRAM based trackers is encouraging, there are still potential of attacks by leveraging cache side channel attacks and probabilistic nature of swaps. We therefore, investigated the design further and modified where needed.

4.3 Tracker and Indirection Design Enhancement

In the previous section, it was observed that despite enhancing the design with a better tracker, the overall memory overhead didn't see much changes. The magnitude of the overall memory overhead is determined by the row-indirection table employed. Therefore, the project was modified to focus on reducing the data structure complication, in terms of maintaining integrated data structures for tracking the rows and then to keep a pointer for identifying swapped row pairs.

One way of achieving reduced storage space is to combine the row-indirection table with the counter and replacing the pair of swap operations with per row swapped address. This led to addition of extra entry equal to the size of one row address.

This led to increase in the overall lookup because now memory has to find the complimentary swap pair through constant look-up in the same table while tracking is in progress. It was also observed that size of entries per bank, ended up with a no reduction in memory overhead since the decrease in the size of one entry is still not enough to compensate for the fact that despite a decreased threshold of $\text{Trh}/6$, the total number of entries inside the tracker is much higher compared to total entries needed for just swapped pairs.

4.4 Accelerated Swap Operations

To reduce the entries inside the indirection table, the design proposal for the accelerated swap operation involved dividing the memory system into multiple swap groups, each containing a set of DRAM banks that can be swapped simultaneously. To implement the accelerated swap operation, we propose a hybrid row swap algorithm that combines the use of the row indirection table and bank swap counters with the parallelism-based approach. The hybrid row swap algorithm works as follows: when a Row Hammer attack is detected, the Swap Group Manager selects two swap

groups that contain the affected row and schedules the swap operations.

This design scheme was not effective due to swap group scope limiting the overall random address space, causing it further easy to probabilistic determination of the swapped pair. Due to above design explorations, it was observed that the major concerns lies in the randomized nature of mitigation. Therefore, the goal is to eliminate the randomized nature of mitigation. This exploration led to leveraging row quarantine region scheme proposed by AQUA. [4]

5 PROPOSED DESIGN

The RRS technique aims to mitigate Row Hammer attacks by using an aggressor-focused scheme by adding an indirection table. However, the technique still relies on randomization, which may not be effective against sophisticated attacks.

By evaluating the approaches mentioned in the proposed design section, it is concluded that randomization based approach is not as effective to be called a fool-proof design. Therefore the next proposal involves reducing the dependencies on buffers to keep the mapping for Row indirection table, and adding a quarantine region to create isolation for row-migration, replacing randomization scheme.

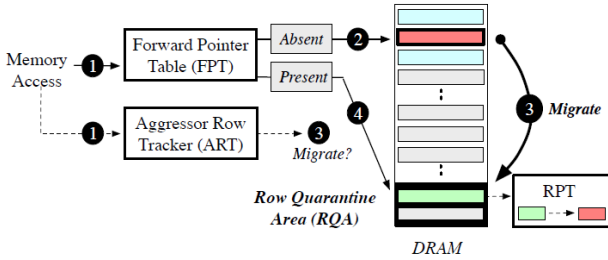


Fig. 3. AQUA design using FPT and RPT alongside the row tracker.

The proposal is inspired by the AQUA design that presents a forward pointer and reverse pointer look-up tables. By observing the previous designs and experiments, it is observed that the tracker when combined with Hydra, which is a SRAM-based tracker, the DRAM memory overhead is improved and also the overall data movement. [4]

The indirection table has now been replaced by a pointer scheme where there is no need to maintain per bank structures. This has seen a reduced memory overhead as well as decrease in total migrations across memory. Now, with Hydra tracker, the total overhead of main memory storage doesn't have to maintain a per bank structure that is increasing in size with threshold, further replicating structures to number of banks. There is also a drop in the overall number of migrations across memory when it is compared with not just RRS but also RRS Hydra modifications, due to a division in lookup where the row-count table is read when group count entries are exploited first.

6 RESULTS

Looking at the results from approach 1, there was no significant change in the performance of the design since even with a hybrid tracker, the design was still using a global mapping table to store row indirection information, Fig. 4. A result of comparison between the designs with and without SRAM trackers shown in Fig. 5 highlights that as

expected the hydra tracker due to accuracy will incur lesser row-migrations but it is not significant enough to overlook the memory overhead of the bulky row-indirection table in each bank.




Fig. 4. Comparison between the normalized cycles RRS MGries and RRS Hydra for benchmarks at Trh of 1024.

Scheme	Threshold	Migrations
rrs	1024	752665
rrs_hy	1024	26790
rrs	2048	427515
rrs_hy	2048	13312

Fig. 5. Comparison between the row-migration in RRS MGries tracker and RRS Hydra design (bzip).

Furthermore, the approach 2 showed no overall benefits on memory overhead by merging the row indirection table with the tracker per bank and adding pointers for isolation. The slight gain in one entry element doesn't recover since it is adding an extra row address size element for tracking that row, even if the threshold hasn't reached.

The quantitative analysis of the tracker and row-indirection table shows that indirection table is incurring order of magnitude greater overhead than tracker. When the tracker is changed to Hydra, the memory overhead caused by RIT is still present, and the number of entries increasing double folds as threshold decreases, Fig. 6.




Fig. 6. The number of entries in RIT with threshold.

The qualitative and quantitative observations from these designs showed no significant performance changes, around 1%. The experiment done on RRS design, by reducing the data structure redundancy despite using Hydra tracker,

hasn't removed probabilistic nature of the design. Therefore, it cannot be reliable for future use.

The final design proposal presents AQUA with hydra tracker, where AQUA has shown better performance in terms of number of cycles as well as reduced number of row migrations, Fig 7.




Fig. 7. Normalized cycles in AQUA implementation compared to RRS with MGries and Hydra tracker.




Fig. 8. Comparison between the row-migration in SRS versus AQUA with hydra.

The non-random row-migration in AQUA has shown considerable improvement in the number of row-migrations, and that has improved further with Hydra tracking, Fig 8. This is also indicative of the fact that lesser number of row-migrations will incur less delay due to reduced memory access for the row-migrations. In this, it was observed that there was a difference in characteristics around benchmarks, when using sufficient number of operations to exploit the RQA data structure with forward and reverse pointer tables, which was attained clearly in bzip. When comparing traces, on bzip it gives a reduction on row-migrations by 7.5 times and on cactus traces by 3 times.

It is also important to note that overall memory overhead has reduced since the per bank RIT has been replaced by a RQA table. Additionally, now because the order of magnitude, in terms of memory storage overhead between the tracker and RQA structure is comparable, the Hydra tracker has incurred 6 times memory overhead reduction.

Moreover, the AQUA Hydra design has also shown some minor overheads. In terms of total memory activation, some benchmarks like bwaves and bzip have shown minor increase compared to victim-focused mitigation. When it comes to total system cycles, benchmarks like cactus and mcf have shown significant increased in cycle for Hydra

extension compared to AQUA design using MGries, Fig 9. It therefore is a matter of application where the gains and overheads vary, however, overall the memory storage gains and reduction in migrations are significant.




Fig. 9. Comparison between the normalized cycles at 1k threshold in AQUA MGries and AQUA Hydra.

7 CONCLUSION

In this project, we proposed an extension to the row-swap technique for mitigating row-hammer attacks in memory systems. Our proposed design aims to address the limitations of the RRS technique and improve its performance and scalability without requiring any additional memory overhead.

Unlike RRS and SRS, AQUA does not rely on randomization, which is a potential security advantage. By integrating the Hydra tracker and modifying the design, we improved the efficiency of design. Our experiments showed that AQUA with non-random row-migration outperformed RRS and SRS in terms of the number of row migrations, memory storage overhead and security. Isolation-based techniques like AQUA have the potential to provide better security than randomization-based techniques for row hammer mitigation.

The future scope of the research indicated potential exploration on the redundancies in the data structures FPT and RPT. The total number of memory access can be lowered down in the AQUA design by flagging off the quarantine addresses in the tracker to point to the actual row address. Furthermore, the challenge lies in evaluation on benchmarks to be able to exploit more tendencies on design by looking at the actual assembly code or instruction sequence in place.

REFERENCES

- [1] M. Qureshi, A. Rohan, G. Saileshwar, and P. J. Nair, "Hydra: Enabling low-overhead mitigation of row-hammer at ultra-low thresholds via hybrid tracking," in *Proceedings of the 49th Annual International Symposium on Computer Architecture*, ser. ISCA '22. New York, NY, USA: Association for Computing Machinery, 2022, p. 699–710. [Online]. Available: <https://doi.org/10.1145/3470496.3527421>
- [2] G. Saileshwar and P. J. Nair, "Scalable and secure row-swap: Efficient and safe row hammer mitigation in memory systems," *arXiv:2212.12613*, 2023.
- [3] G. Saileshwar, B. Wang, M. Qureshi, and P. J. Nair, "Randomized row-swap: Mitigating row hammer by breaking spatial correlation between aggressor and victim rows," in *Proceedings of the 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, ser. ASPLOS '22. New York, NY, USA: Association for Computing Machinery, 2022, p. 1056–1069. [Online]. Available: <https://doi.org/10.1145/3503222.3507716>
- [4] A. Saxena, G. Saileshwar, P. J. Nair, and M. Qureshi, "Aqua: Scalable rowhammer mitigation by quarantining aggressor rows at runtime," in *2022 55th IEEE/ACM International Symposium on Microarchitecture (MICRO)*, 2022, pp. 108–123.